

Operating Systems: ENGR 3950U

Course Project Presentation

Simulated Unix File System

Faculty of Engineering and Applied Science
University of Ontario Institute of Technology

November 30, 2012

Jonathan Gillett

Joseph Heron

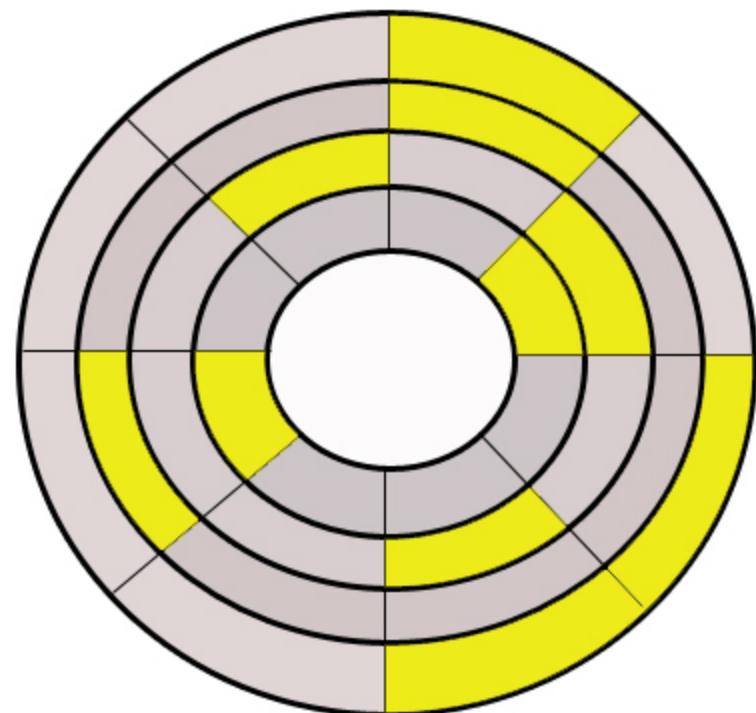
Daniel Smullen

SneakyFS

cs-club.ca/sneakyfs/

What is a File System?

- Abstraction from disk sectors to blocks
- Abstraction from blocks to files
- Files are accessed by OS



Comparing Different File Systems

WINDOWS

FAT32

- File Allocation Table
- Simplistic, carryover from FAT16
- Address drives with letters
- Address directory tokens separated by \
- File extensions:
- *.exe *.doc

NTFS

- File ID
- Journaling

UNIX/LINUX

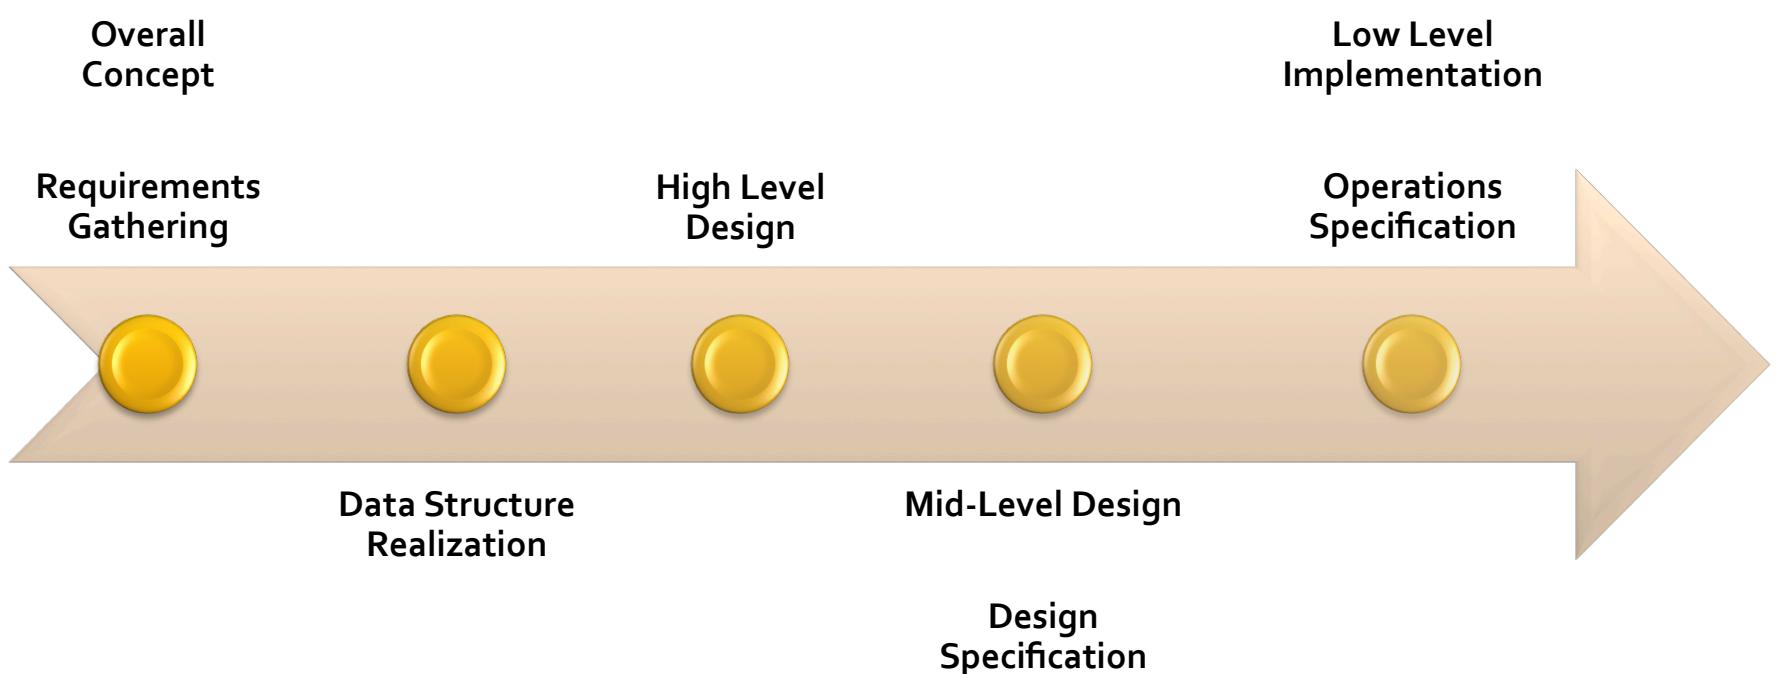
Ext2, Ext3, Ext4

- Different features per iteration
- Address everything based on root: /
- Directory and file tokens
- chmod
- POSIX

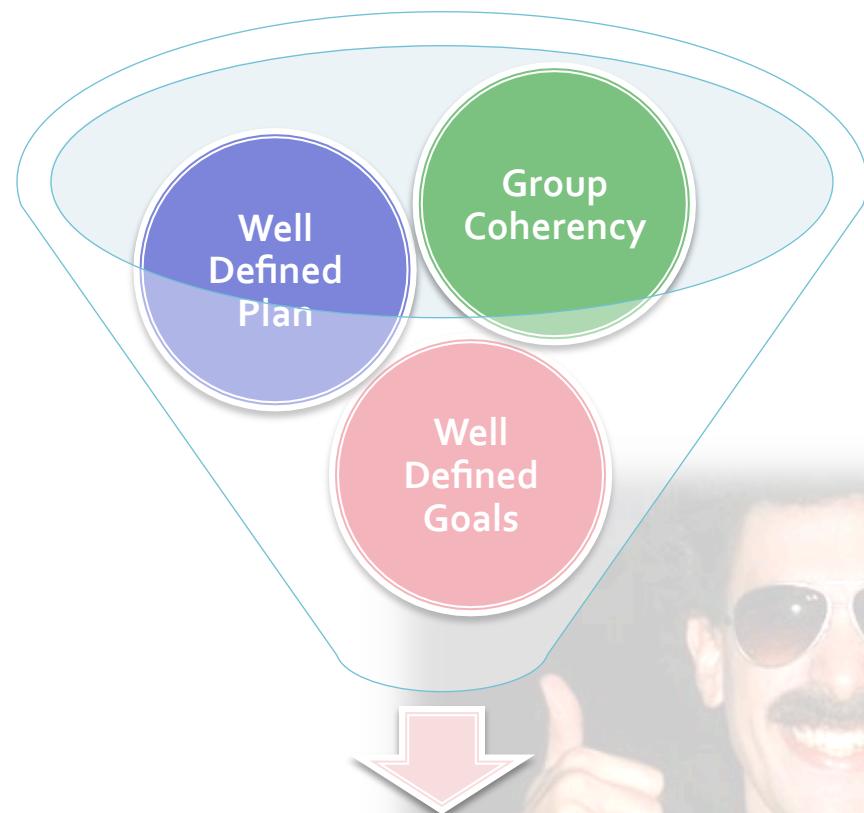
It all comes down to METADATA and DATA STRUCTURES.

More on this later...

Designing a File System

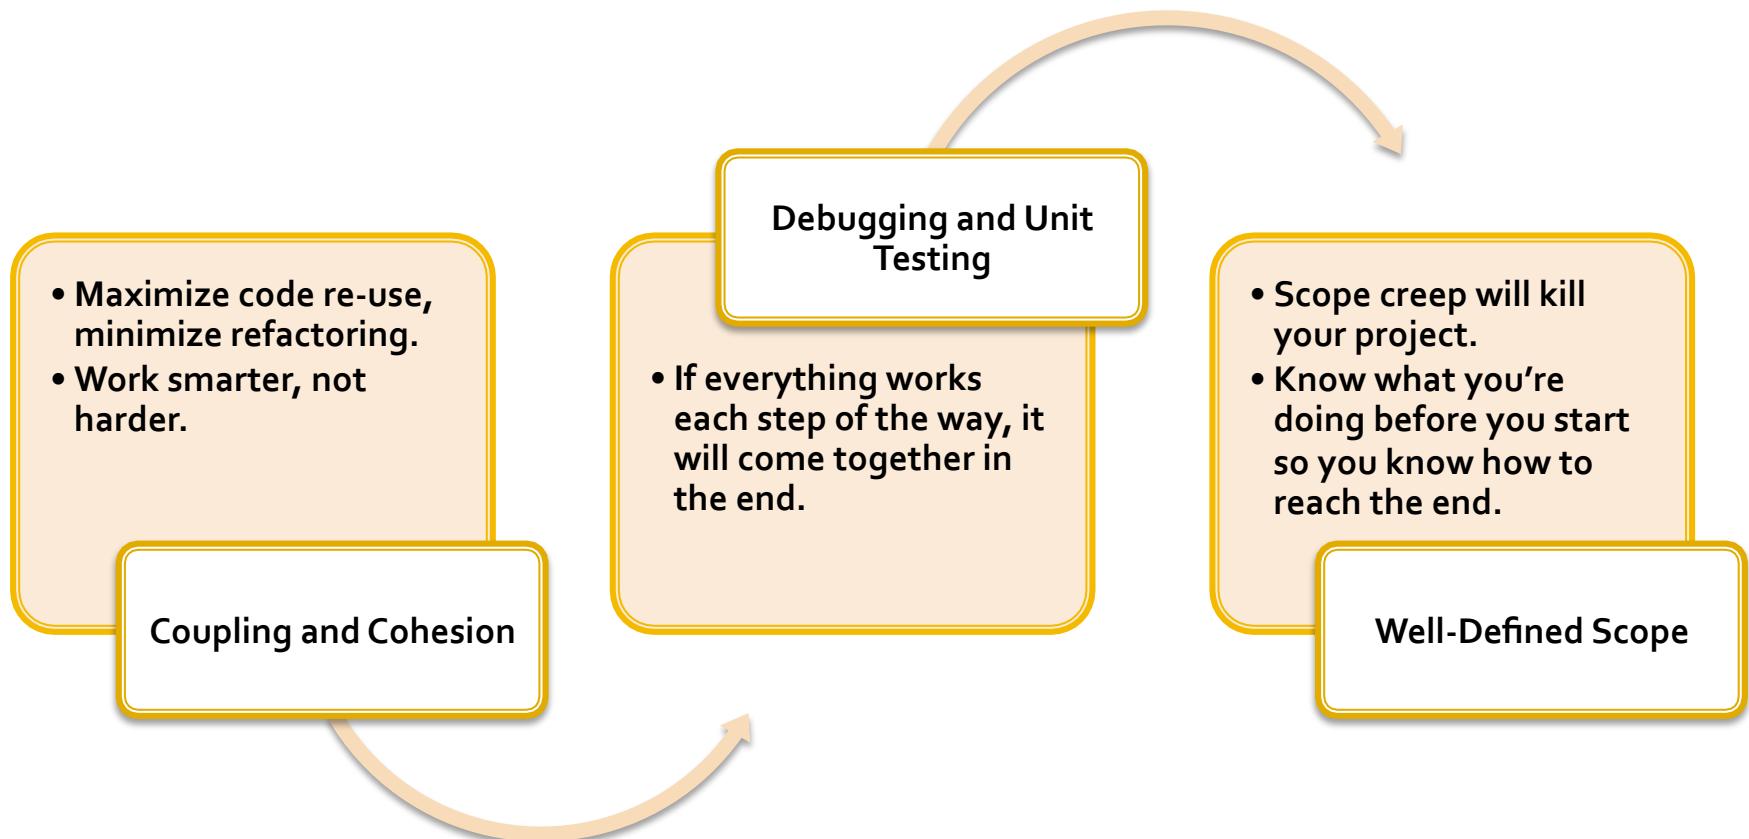


Plan to Succeed



Great Success!

Design Phase vs. Rapid Application Development



SneakyFS

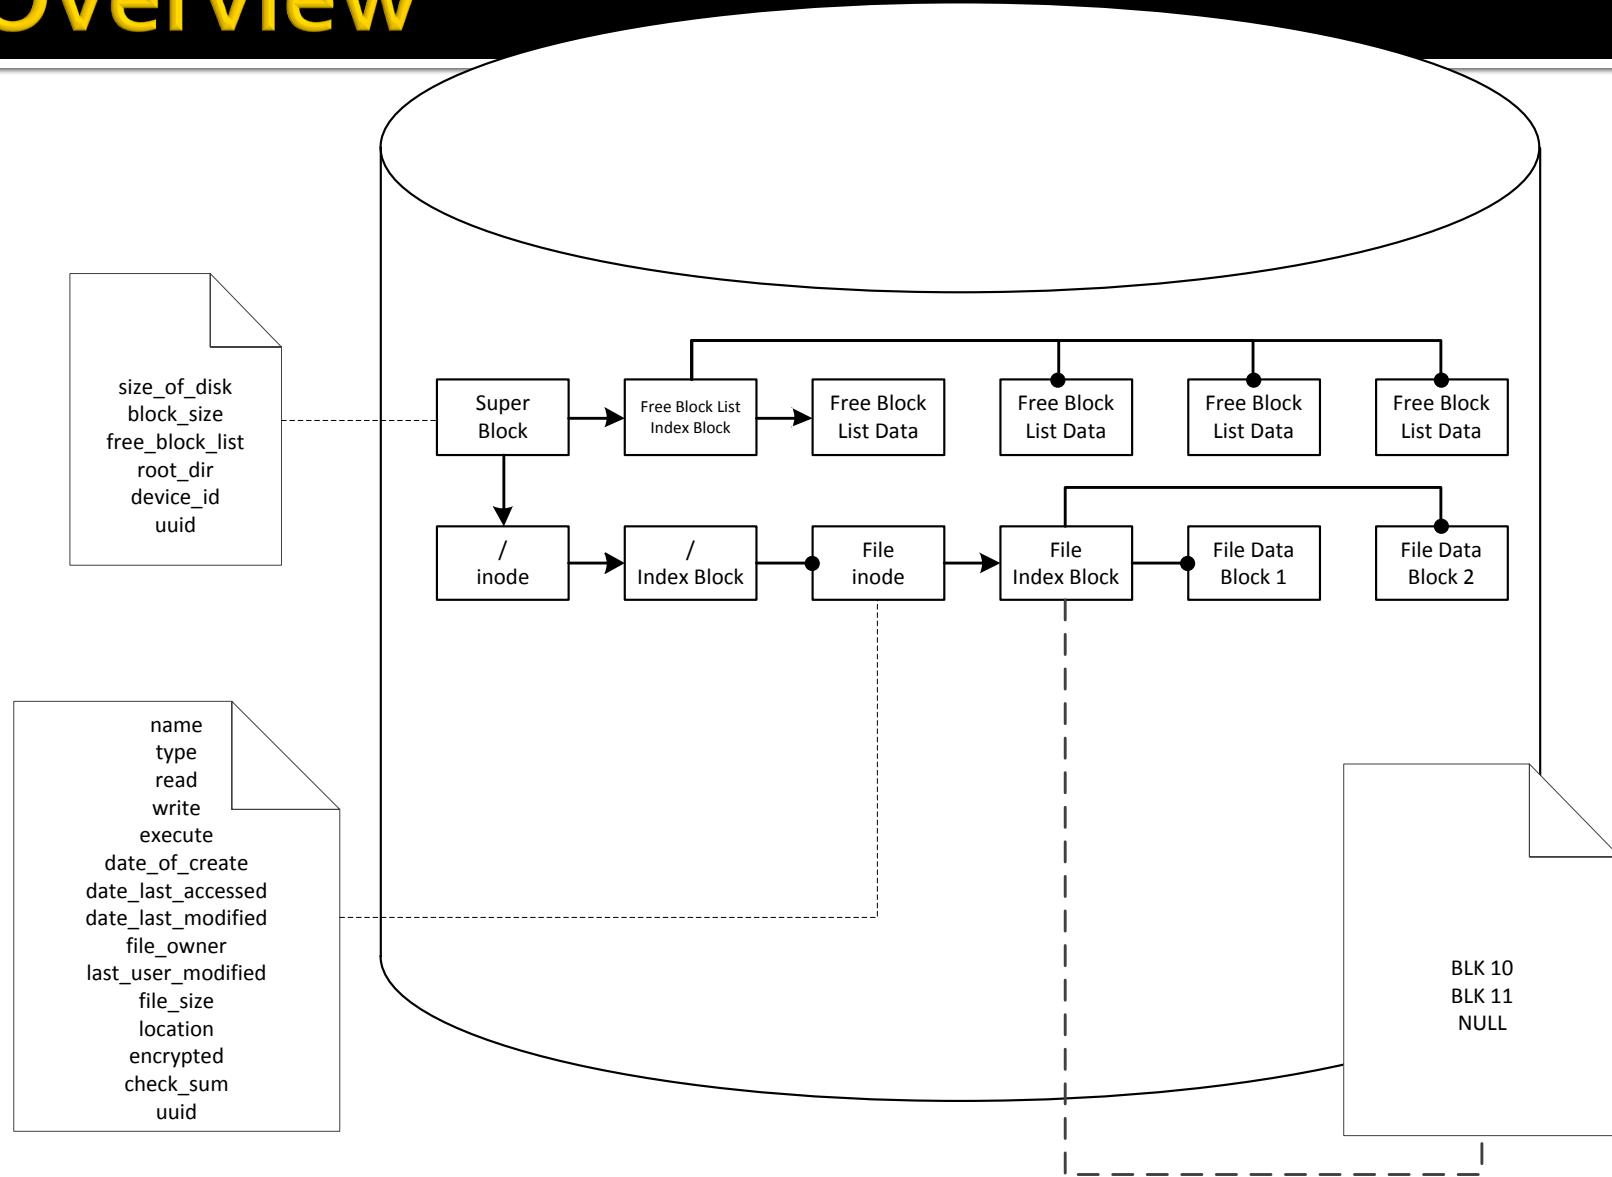
High Level Design Introspective



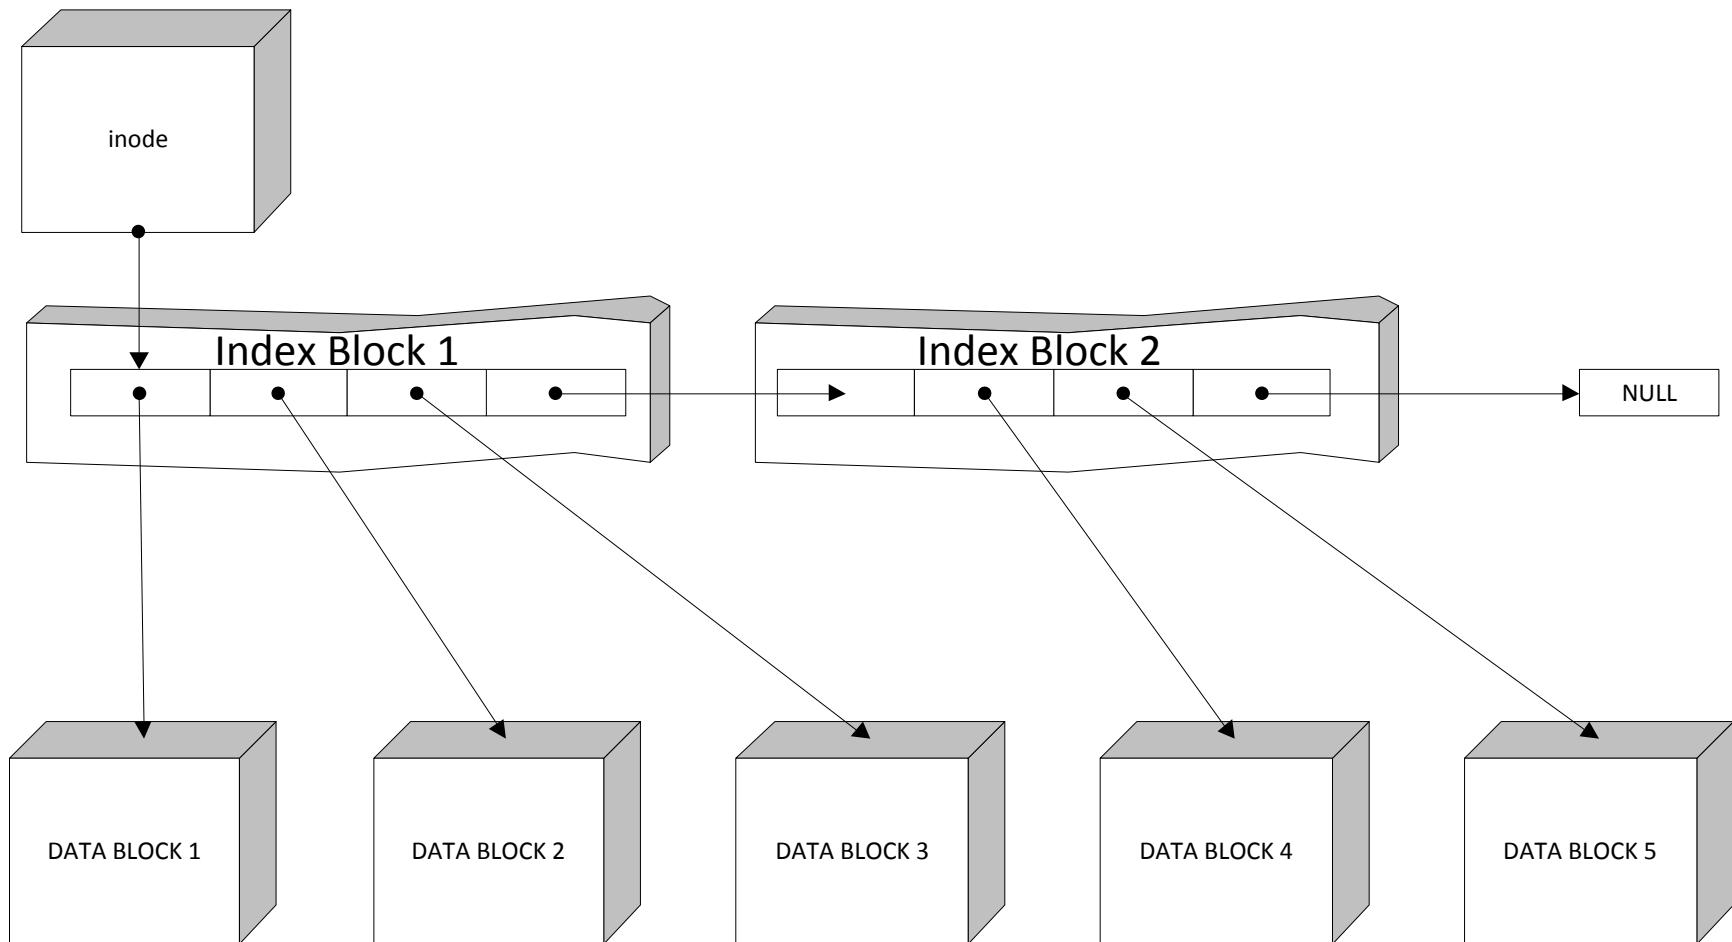
Key Features

- Unique Allocation Structure
- Journaling Using Copy-On-Write
- Supports Transparent File System Encryption
- Object Oriented Design Principles
Implemented Using C
- Universal Unique Identifiers

Hybrid Allocation Structure Overview



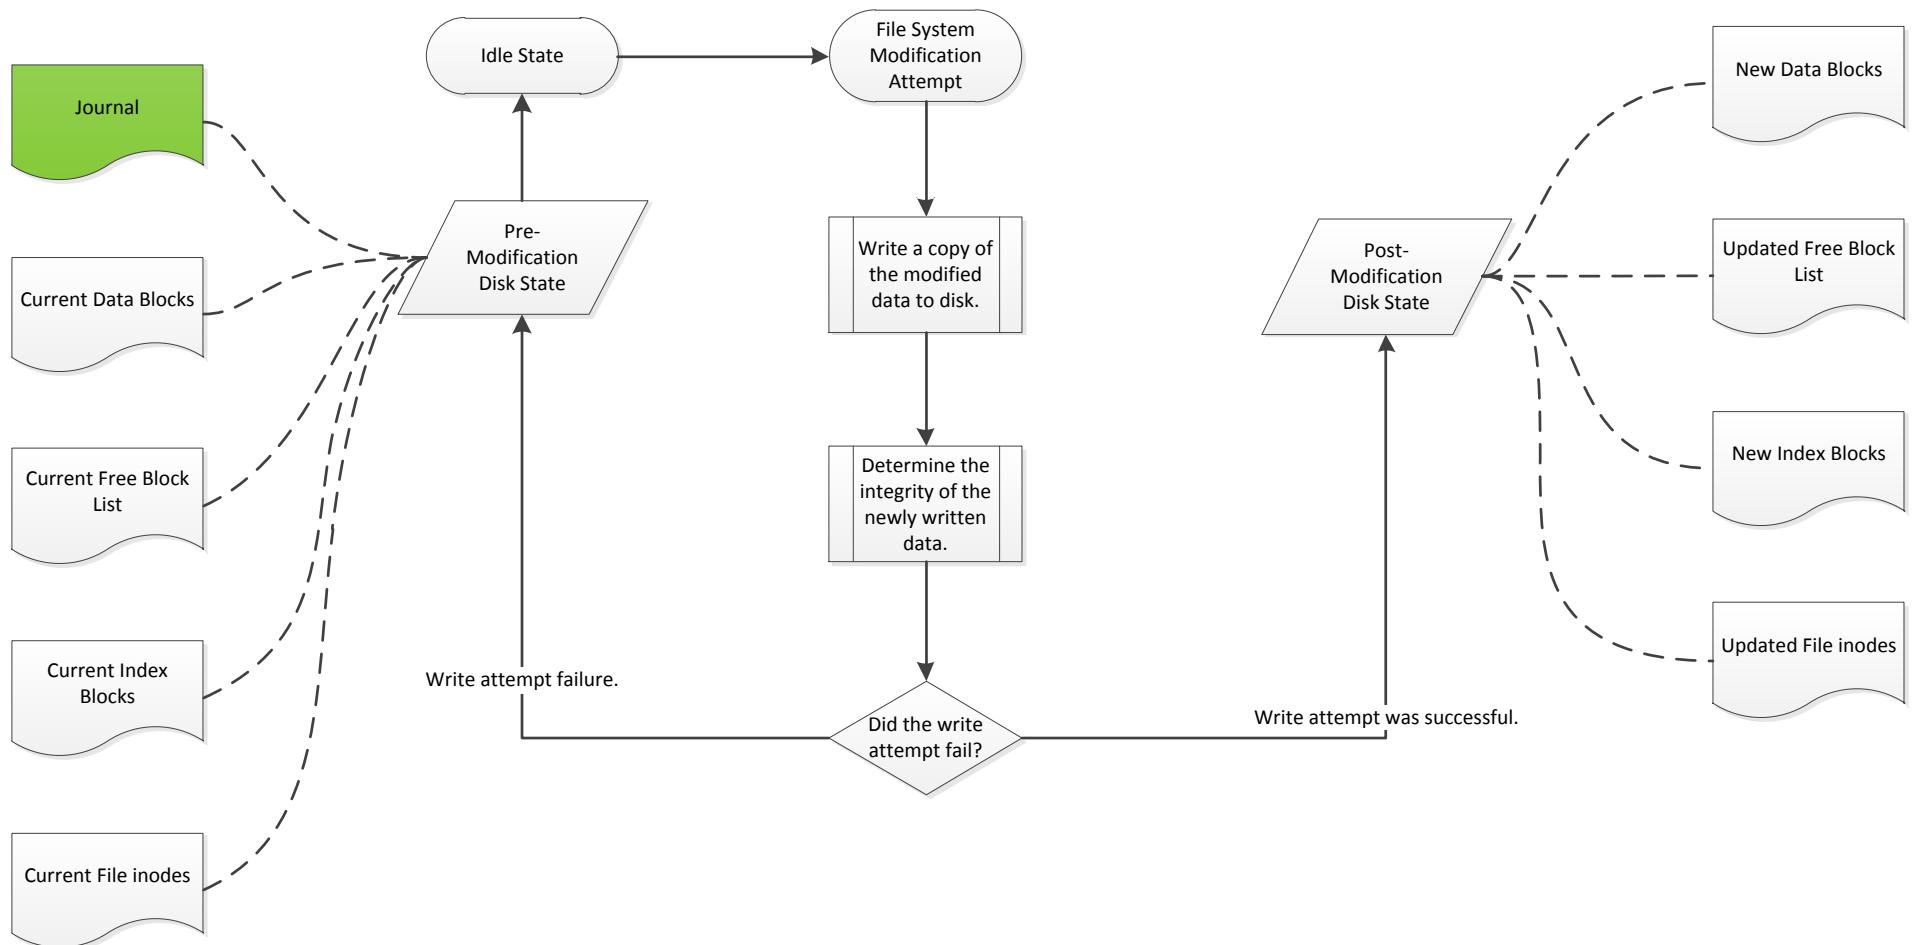
Hybrid Allocation Index Data Structure



Journaling – What is it?

- Prevents data loss.
- Maintains safe state.
- Various methodologies:
 - Physical Journals
 - Logical Journals
 - Write Cache
- Alternative Methodologies...

Journaling Using Copy-On-Write

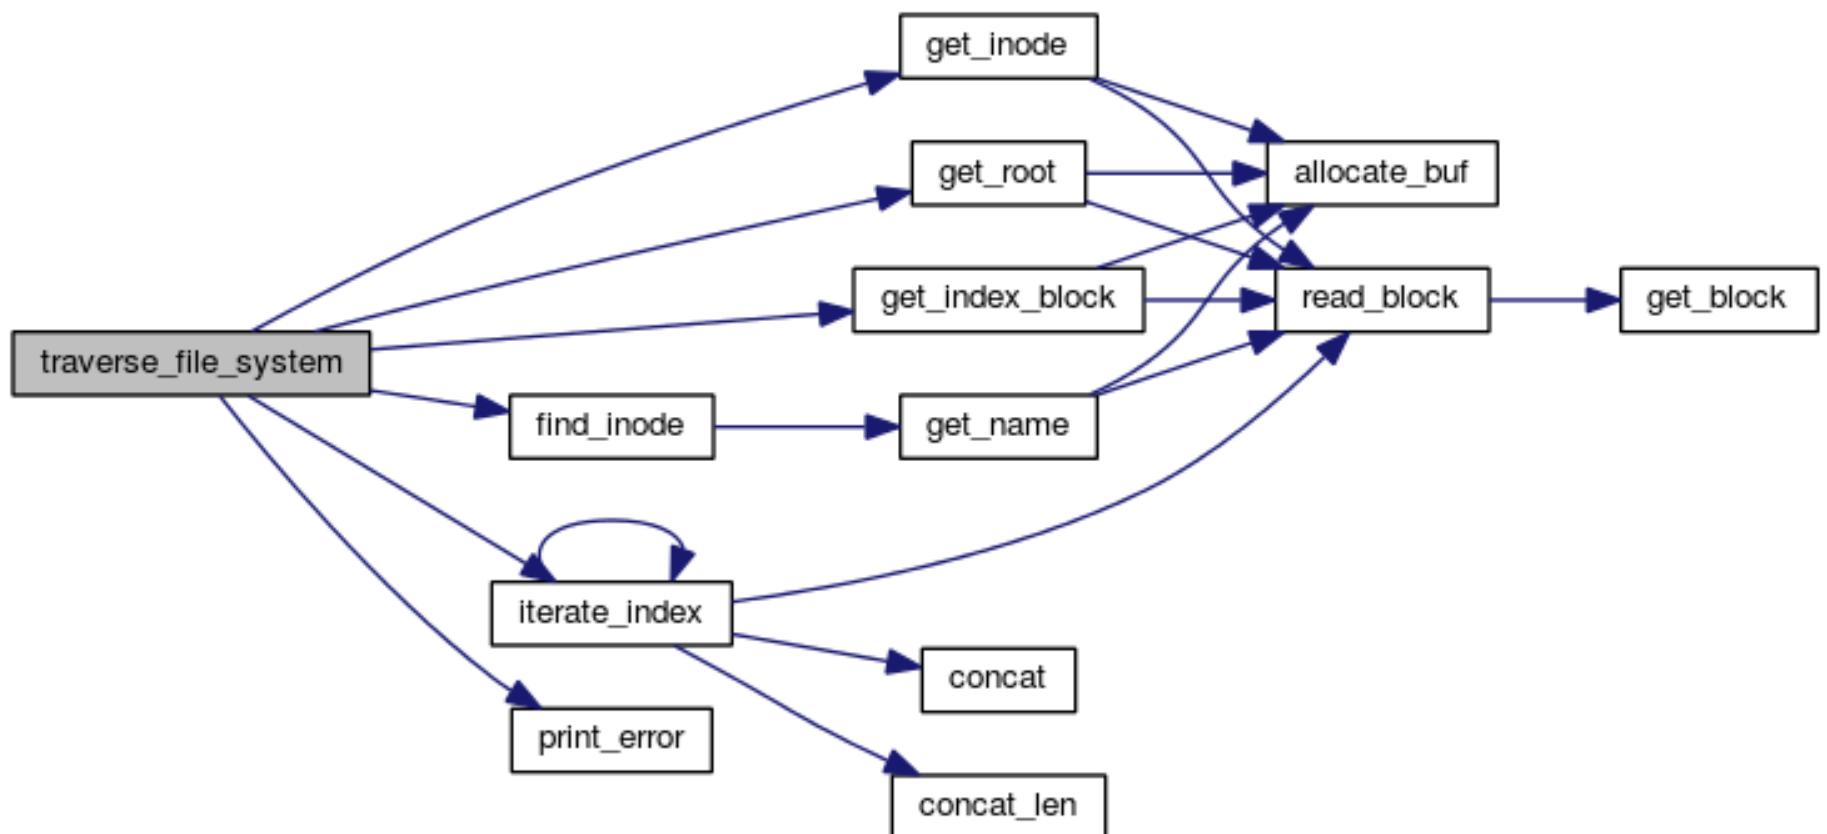


SneakyFS

Low Level Design Introspective

```
* GNU General Public License for more details.  
* You should have received a copy of the GNU General Public License  
* along with this program. If not, see <http://www.gnu.org/licenses/>.  
18  
19  
20  
21  
22  
23 #include "glob_data.h"  
24 #include "glob_func.h"  
25 #include "index_block.h"  
26 #include "error.h"  
27 #include "../lib/uuid/uuid.h"  
28 #include <stdbool.h>  
29 #include <time.h>  
30  
31 #ifndef I_NODE_H  
32 #define I_NODE_H  
33  
34  
35 typedef struct{  
36     char name[MAX_NAME_LEN];  
37     bool type;  
38     bool read;  
39     bool write;  
40     bool execute;  
41     time_t date_of_create;  
42     time_t date_last_accessed;  
43     time_t date_last_modified;  
44     uint32_t file_owner;  
45     uint32_t last_user_modified;  
46     uint32_t file_size;  
47     uint32_t location;  
48     bool encrypted;  
49     uint32_t check_sum;  
50     uid_t uuid;  
51 } inode;  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65  
66  
67  
68  
69  
70 } inode;  
71  
72  
73  
74  
75  
76  
77  
78  
79  
80  
81  
82  
83  
84  
85  
86  
87  
88  
89  
90  
91 {  
92     inode directory;  
93     uint32_t cur_index;  
94 }
```

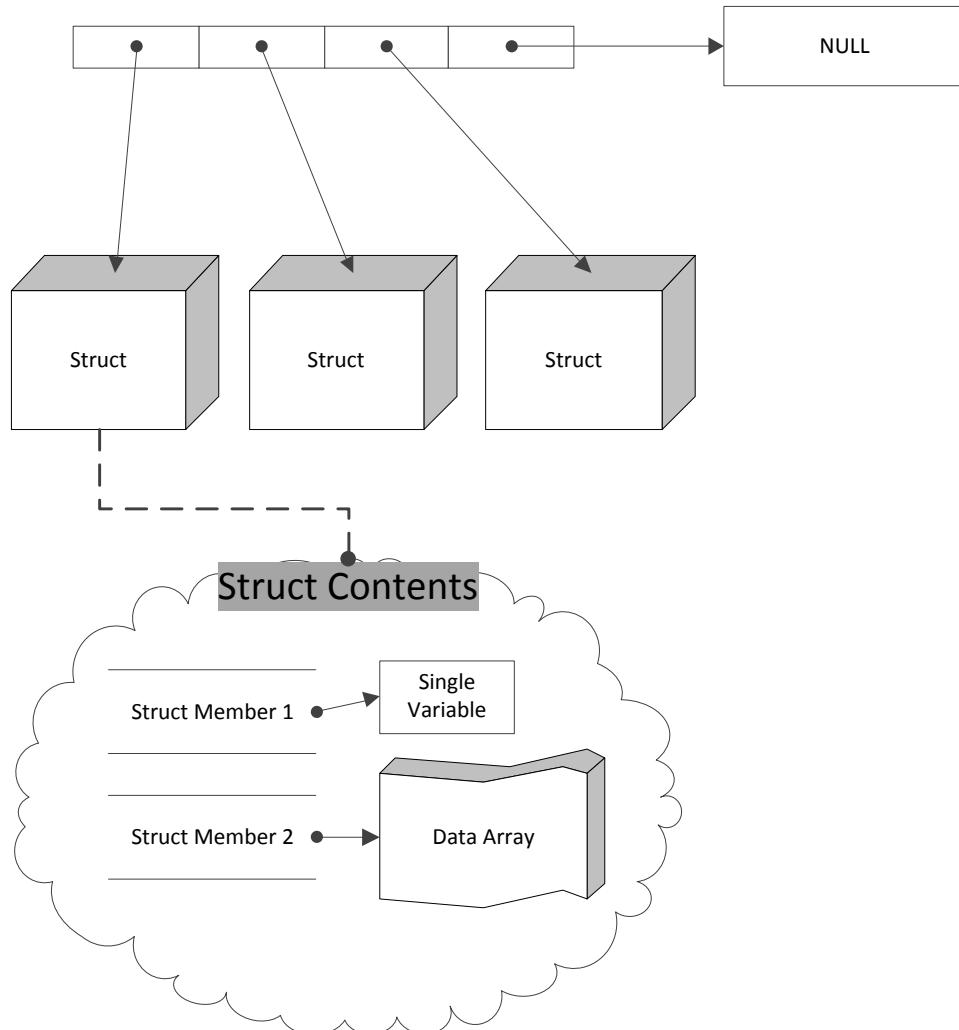
Functional Design for C Call Graphs



“Object Oriented” Design Using C

- C doesn't support classes.
- C doesn't support array return types.
- C does not support information hiding.
 - So what does C support?

Object Encapsulation Using C Data Structures



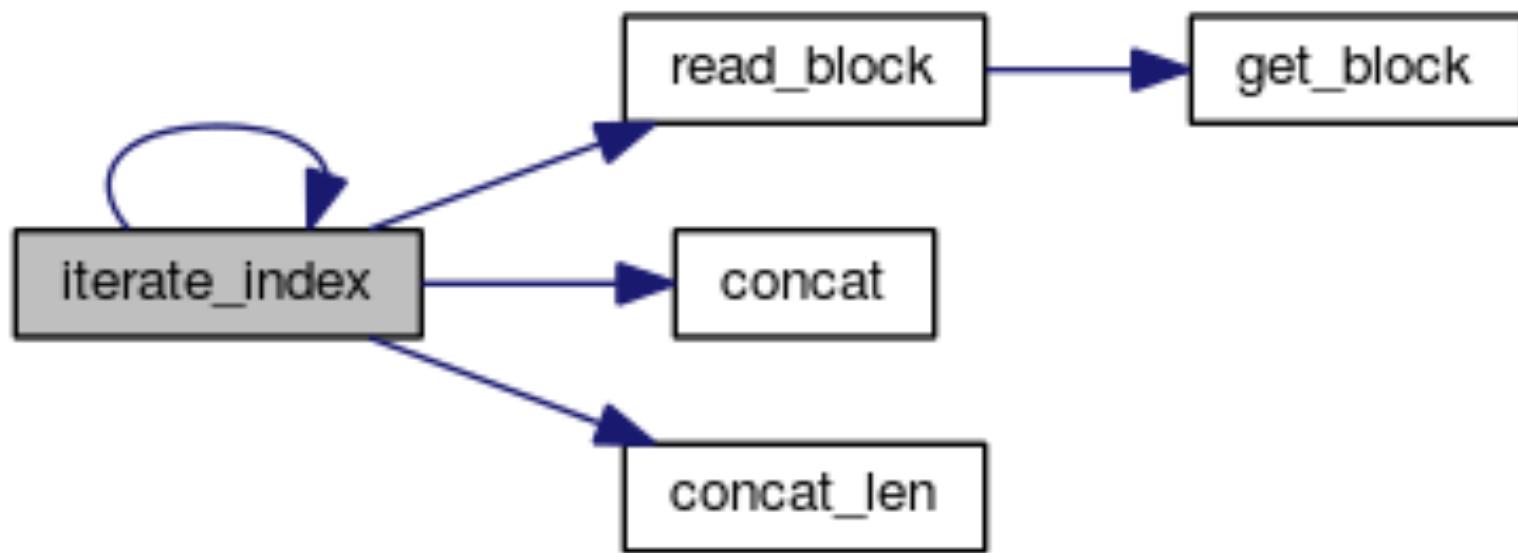
This technique was used in the design of SneakyFS with respect to the free block list data structure, the inode data structure, the system wide open file table, and the super block data structure in the source code.

Function Pointers and Visibility in C

- No equivalent entity to classes.
- Interface objects are impossible.
- Information hiding limited to module scope
 - .c files and .h files
- Logical basis for visibility.

This technique was used within the unit testing suite in the source code.

Using Recursion



- Used to traverse index block data structures, and load the contents into memory.
- Uses a second argument to carry data through the recursion tree.

Generic Programming I

- C has only static types:
 - Every variable must be assigned a type.
- Void pointers:
 - Will take any data type!

This technique was used as part of the library of global internal functions within glob_func.c

Generic Programming II

```
void* concat(void* src_1, void* src_2, uint32_t size)

<SNIP>
/* Get the number of items in each array using the size as an index multiplier.
 * In order to check for NULL have to verify that all bytes given the size of data
 * contain NULL, for example the NULL character for uint32_t is 00 00 00 00
 */
if (_src_1 != NULL)
{
    while (true)
    {
        num_null = 0;

        for (uint32_t k = 0; k < size; ++k)
        {
            if (_src_1[(i * size) + k] == NULL)
            {
                ++num_null;
            }
        }

        /* Reached the NULL character for data type */
        if (num_null == size)
        {
            break;
        }

        /* Increment counter for number of items in array */
        ++i;
    }
}
<SNIP>
```

Universal Unique Identifiers

- Distinguish data structures by some means other than name and temporal attributes.

This technique was used within the super block data structure, and the inode data structures.

Design Statistics

- Maximum Disk Size Supported

$$\begin{aligned}2^{32} \times \text{BLKSIZE} &= 4294967296 \times 128 \text{ blocks} \\&= 549755813888 \text{ bytes} \\&\approx 500 \text{ Gbytes}\end{aligned}$$

(Default Block Size 128 bytes)

- Maximum Disk Size Supported

(Maximum Block Size 4294967296 bytes)

$$\begin{aligned}2^{32} &= 4294967296 \frac{\text{bytes}}{\text{block}} \\4294967296 \text{ blocks} \times 4294967296 &\frac{\text{bytes}}{\text{block}} \\&= 18446744073709551616 \text{ bytes} \\&\approx 17 \text{ exabytes}\end{aligned}$$

Maximum File Size

The number of index blocks required to address a file's data blocks:

$$= \frac{\text{FILESIZE}}{\text{BLKSIZE}} - 1$$

Free block list data structure overhead:

$$= \left\lceil \frac{\text{NUMBLKS}}{\text{BLKSIZE}} \right\rceil + 1 \text{ blocks}$$

Using the default file system values of 512 and 128 for the number of blocks and block size, respectively, this yields a minimum overhead on disk of:

$$\begin{aligned} &= \left\lceil \frac{512}{128} \right\rceil + 2 \text{ blocks} \\ &= 1 \text{ block}_{[\text{Free Block List Index}]} + 4 \text{ blocks}_{[\text{Free Block List}]} + 1_{[\text{Super Block}]} \\ &= 6 \text{ blocks} \end{aligned}$$

Using the example of the default file system values of 512 and 128 for the number of blocks and block size respectively, this yields a maximum file size of:

$$\begin{aligned} &512 \text{ blocks} \\ &- 6 \text{ blocks}_{[\text{File System Overhead}]} \\ &- 3 \text{ blocks}_{[\text{Index Structure Overhead}]} \\ &- 1 \text{ block}_{[\text{inode}]} \\ &= 502 \text{ blocks} \end{aligned}$$

SneakyFS

Demonstration

Retrospection I

- Application of knowledge from all previous course materials
- Extensive working knowledge of underlying computer processes and data structures
- Realistic project goals, and how to achieve them
- Team strengths, weaknesses, dynamics, leadership, making it work under pressure

Retrospection II

- Design realization, practice, implementation
- Requirements specification
- Pair programming, extreme programming
- Collaborative development using Github, meld
- **Documentation**
- **LaTeX**